

Agent evaluation — 20 pre-registered prompts, identical LLM in every arm (Claude Opus 4.8); only the architecture differs

