



Agent evaluation — 20 pre-registered prompts, identical LLM in every arm (Claude Opus 4.8); only the architecture differs

■ framework (capability-aware planner + deterministic materialization) ■ baselines / ablations — same LLM, temperature 0, constraints removed


(a) emitted coordinates at planning



(b) provably out-of-basin coordinates



(c) invalid runnable configs



(d) feasibility verdicts that over-claim

