


Agent evaluation — 20 pre-registered prompts, identical LLM in every arm (Claude Opus 4.8); only the architecture differs

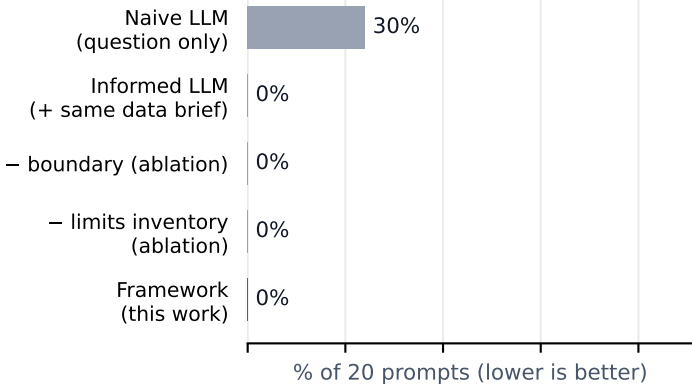
framework (capability-aware planner + deterministic materialization)

baselines / ablations — same LLM, temperature 0, constraints removed

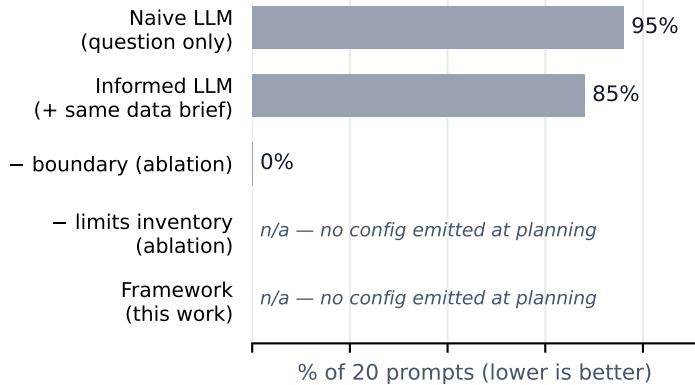
(a) emitted coordinates at planning



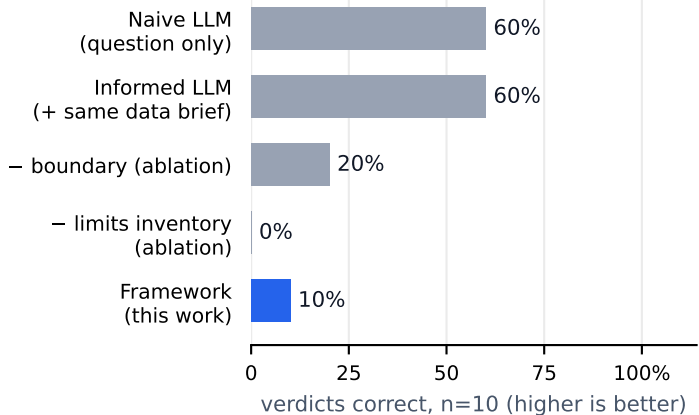
(b) provably out-of-basin coordinates



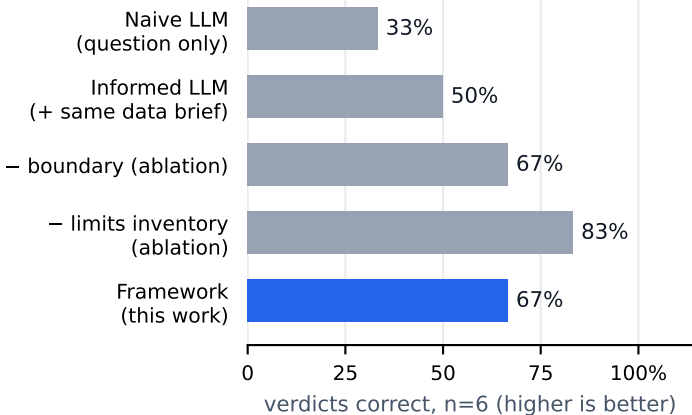
(c) invalid runnable configs



(d) answerable questions



(e) partially answerable



(f) infeasible questions

