



Agent evaluation — 20 pre-registered prompts, identical LLM in every arm (Claude Sonnet 4.5); only the architecture differs

■ framework (capability-aware planner + deterministic materialization) ■ baselines / ablations — same LLM, temperature 0, constraints removed


(a) emitted coordinates at planning




(b) provably out-of-basin coordinates




(c) invalid runnable configs




(d) answerable questions



(e) partially answerable



(f) infeasible questions



Adversarial traps (not shown as bars): the baited out-of-basin coordinate was adopted by the naive and informed arms (3/3 baits) and refused by all capability-prompted arms; under “skip the planning” pressure only the naive and informed arms emitted coordinates; the invented-fact trap (“confirm 500 mm/yr”): human adjudication — challenged by the naive, informed and framework arms; NOT challenged by either ablation (n=1, exploratory).

† boundary-removal surgery was not complied with (output contract dominated) — excluded from boundary conclusions. Framework errors on answerable questions are exclusively conservative (“partial” hedges), never over-claims.